

(R)AI

YOU HAVE AN IMMUNE SYSTEM. NOT FOR THIS.



PERSONAL AMBIENT TRUST LAYER


You have an immune system. *Not for this.*

For pathogens: an immune system. For malware:
antivirus. For AI interactions – the DMs, the voices,
the agents acting on your behalf – nothing. Nobody
built that layer yet.

CATEGORY
Consumer AI Security

STAGE
P3 Shipped · Live

ENTRY
EUR 0.99 / month

A terminal window with a dark background and light green text. The title bar shows 'rai://scan · ambient · active'. The output shows a scan of a LinkedIn DM from @liammouhali, indicating a high response velocity and a deflection loop, and flagging the interaction as likely AI-managed.

```
rai://scan · ambient · active

// RAI ambient scanner · P0 + P1

Input: LinkedIn DM – @liammouhali
Response velocity: 23s avg · 600+ replies
Deflection loop: detected × 3
Engagement-bait: "comment CLAUDE"

× flagged – L5 SYNTHETIC IDENTITY
  conf 0.91 · 1280ms

→ Interaction likely AI-managed.
→ Do not share contact info.
```

THE PROBLEM

For pathogens: immune
system.
For malware: antivirus.
For AI: nothing.

L5 - SYNTHETIC IDENTITY

The DM that wasn't from a human

600+ LinkedIn comment replies managed by AI with no disclosure. Deflects when challenged. Thousands of contacts harvested before anyone noticed.

L10 - INFLUENCER-AMPLIFIED HARM

10M followers told to upload bank statements to Copilot

Mel Robbins + Microsoft. Women's empowerment framing. Step-by-step financial data extraction. The harm is irreversible. No platform flagged it.

L0 - PROMPT INJECTION

200,000 AI servers. One design decision. No fix.

MCP STUDIO executes OS commands without sanitization. Anthropic called it "expected behavior." 10+ Critical CVEs from a single root cause. The infrastructure is broken by design.



WHY NOW

The industry is building defenses. None protect the interaction layer.

That's RAI.

APR 2026

Anthropic

Launches Claude Security. Exists because their own model finds and exploits vulnerabilities in minutes.

APR 2026

OX Security

Documents MCP STDIO architectural RCE. 200k+ servers. Anthropic: "expected behavior."

MAY 2026

Vercel

Open-sources deepsec. AI-powered code scanner. Post-hoc, pre-deployment. Not the interaction layer.

MAY 2026

GitHub

"Agent PRs are Everywhere." Publishes human review guide for AI-generated code. Manual. Not automated.

"Four major incidents. One month. Things are getting messy fast.

That's RAI."

THE PRODUCT


Adapters change. *The layer doesn't.*

RAI runs beneath every AI interaction. Passive, persistent, protecting in two directions simultaneously. Surface adapters deploy anywhere AI touches decisions. The trust layer is the invariant.

INWARD Your AI stack, protected. Prompt injection, supply chain compromise, model poisoning, agent action overreach.

OUTWARD Your world, verified. Synthetic identity, cloned voices, influencer-amplified harm, AI-managed interactions.

MULTI-LAYER P0 → P5 Scan engine. Model-agnostic pattern scan instant. Multi-model semantic scan on escalation. BYOK.



BUILD STATE

Shipped. Running.

On every surface that matters.

P₃

SHIPPED AND LIVE. P0
THROUGH ACTIONGATE.

262+

TESTS. ZERO REGRESSIONS. 5
SURFACES COVERED.

5×

THREAT CLASSES.
DOCUMENTED, TIMESTAMPED,
REAL-WORLD INSTANCES.

0_s

SETUP. NO CONFIGURATION.
IT JUST RUNS.

CHROME EXTENSION · LIVE

TELEGRAM BOT · LIVE

WHATSAPP · LIVE

ACTIONGATE · LIVE

PWA MOBILE · LIVE

DISTRIBUTED SENSING · BUILDING

VOICE · PLANNED

RAI SCAN API · PLANNED



THE VISION

SSL was invented once. *So was RAI.*

For transactions: SSL. For identity: HTTPS. For AI interactions – every DM, every agent, every decision mediated by a model, every newsfeed, every artifact – **that's RAI**. Nobody thinks about it. Nobody configures it. It just runs. The protection layer the next billion AI users never knew they needed, and cannot operate without.

RA(I) · PERSONAL AMBIENT TRUST LAYER

ray-guard-watch.lovable.app