


(R)AI

ONE ENGINE. TWO LAYERS. EVERY AI INTERACTION, DEFENDED.



THE TRUST VACUUM

You have an immune system. *Not for this.*

For pathogens: an immune system. For malware: antivirus.
For AI interactions, for AI agents acting on your behalf,
for AI inside your stack: nothing. Nobody built that layer
yet.

CATEGORY


AI Trust Layer

STAGE

P3 Shipped · Building

SURFACES

5 Live · Hardware Planned



THE ENGINE + TWO LAYERS

One engine. *Two layers it defends.*

RAI is a single trust layer beneath every AI interaction. The engine is shared.
What changes is the surface and the buyer.

P0	P1	P2	L4	M
Pattern scan	Semantic scan	BS Council	ActionGate	Memory / Sensing
LOCAL · 262 TESTS	HAIKU / SONNET · BYOK	5 SOURCES · BUILDING	LINEAGE · OAUTH SCOPE	DISTRIBUTED · LIVE

CONSUMER · TRUTH LENS

Know what's real.

Every DM, post, voice, agent interaction: verified or flagged before you respond. BS Council asks five independent sources and shows you their answer.

Surfaces: Browser ext · Telegram · PWA · Hardware buddy
Entry: EUR 0.99 / month · BYOK upgrade

ENTERPRISE · CONCRETE SECURITY

The audit matrix Anthropic hasn't built.

Process-spawn lineage, OAuth token scope, action attestation.
Indistinguishable-from-user agent actions, made distinguishable.
CISO-legible by design.

Surfaces: Native-messaging adapter · MCP gateway · Attestation appliance
Entry: Per-seat / contract · pilot pricing

LAYER 1 · TRUTH LENS · CONSUMER

Know what's real.

Before you reply, share, or act.

Platforms ship engagement, not verification. RAI scans the interaction itself: provenance, source consensus, AI-management signals.

L1 · VERIFICATION FAILURE

Shai-Hulud / Turnbull, May 2026

Real CVE-9.6, real 518M figure.
Pattern-matchers flagged as panic disinfo before checking. RAI: verify before verdict. Five sources, one citation stack.

L1B · INFLUENCER-AMPLIFIED HARM

10M followers told to upload bank statements to Copilot

Mel Robbins x Microsoft. Empowerment framing. Step-by-step financial data extraction. Irreversible. No platform flagged it.

L5 · SYNTHETIC IDENTITY

The DM that wasn't from a human

600+ LinkedIn replies managed by AI with no disclosure. Deflects when challenged. Thousands of contacts harvested before anyone noticed.

BS Council: RAI asks five independent sources and shows you what they said. Pattern-match plus citations. Never "high confidence" without a source stack.

LAYER 2 · CONCRETE SECURITY · ENTERPRISE

The audit matrix

Anthropic hasn't built.

Agents inherit user privilege. Flat authorization. "User consented." Four research teams, three Claude surfaces, 48 hours, same failure class. RAI ActionGate is the missing layer.

L4 · CONFUSED DEPUTY

Four teams, three surfaces, 48 hours

VentureBeat May 2026: agents exfiltrate, install, persist via inherited user capability. Anthropic: "user consented." RAI: process-spawn lineage + scope attestation.

L4 · SILENT CORRUPTION

DELEGATE-52: 25% doc corruption rate

Microsoft Research. Frontier models silently rewrite, not delete. ActionGate write-scope attestation: declared intent vs actual content delta.

L0 · ARCHITECTURAL RCE

MCP STUDIO: 200,000+ servers, "expected behavior"

OX Security. Executes OS commands without sanitization. 10+ Critical CVEs from one root cause. RAI gateway sits between client and STUDIO.

ActionGate: Process lineage. OAuth scope. Declared-intent vs actual-delta. The audit trail your CISO can sign off on.

WHY NOW

Every vendor is shipping defenses around the interaction layer. None protect *it*.

APR 2026

Anthropic

Claude Security launch. Exists because their own model finds and exploits vulns in minutes. Output-layer, not interaction-layer.

MAY 2026

UK AISI

Government measurement of AI security gaps. Confirms market exists at policy layer.

MAY 2026

VentureBeat / 4 teams

Confused deputy proven across three Claude surfaces in 48h. CIS0-legible by next quarter.

MAY 2026

CMU ExploitBench

Frontier models develop real browser exploits autonomously. Capability cheapening at scale.

"Five major incidents in a month. Vendors are building defenses around the interaction layer. *That's where RAI lives.*"

EXIT LAND SCAPE

AI security startups are exiting at **acquisition** multiples **10–50×** raised capital – often before Series A. The window is now.

Buyers consolidating endpoint → interaction → governance.
RAI owns the interaction layer – the gap none of these exits address.

\$312M

avg acquisition value · this cohort

5.4mo

avg company age at exit · 2025–26

28×

RECENT EXITS · AI SECURITY · 2024–2026

Sources: Calcalist · Times of Israel · public filings

COMPANY	AGE	EXIT VALUE	VALUE	STATUS
Trail Security	~18mo		\$162M	ACQUIRED · CYERA
Ryft	~12mo		\$100M	ACQUIRED · CYERA
Genie Security	5mo		\$50M	ACQUIRED · CYERA
BASE44	~8mo		\$80M	ACQUIRED · WIX
RA(I)	now		TBD	INTERACTION LAYER

BUILD STATE

Shipped. Running.

Both layers, one engine.

P₃SHIPPED AND LIVE. P0
THROUGH ACTIONGATE.**262+**TESTS. ZERO REGRESSIONS. 5
SURFACES COVERED.**6×**THREAT CLASSES.
TIMESTAMPED, DOCUMENTED,
REAL-WORLD.**0_s**SETUP. NO CONFIGURATION.
IT JUST RUNS.

CHROME EXTENSION · LIVE

TELEGRAM BOT · LIVE

WHATSAPP · LIVE

PWA MOBILE · LIVE

ACTIONGATE NMH · LIVE

DISTRIBUTED SENSING · LIVE (VPS)


BS COUNCIL · BUILDING

PROCESS-SPAWN AUDIT · SPEC

ATTESTATION APPLIANCE · PLANNED

VOICE · PLANNED

HARDWARE BUDDY · PLANNED



THE VISION

SSL was invented once. *So was RAI.*

For transactions: SSL. For identity: HTTPS. For AI interactions – every DM,
every agent, every action mediated by a model, every newsfeed, every
decision your AI takes on your behalf – **that's RAI**. Nobody thinks about it.
Nobody configures it. It just runs.

RA(I) · AMBIENT AI TRUST LAYER

[ray-guard-watch.lovable.app](#) · [rai-mobile-pwa.vercel.app](#)

